

This represents the default. Through additional configuration, `kubeadm init` can behave differently, such as reusing an existing etcd cluster.

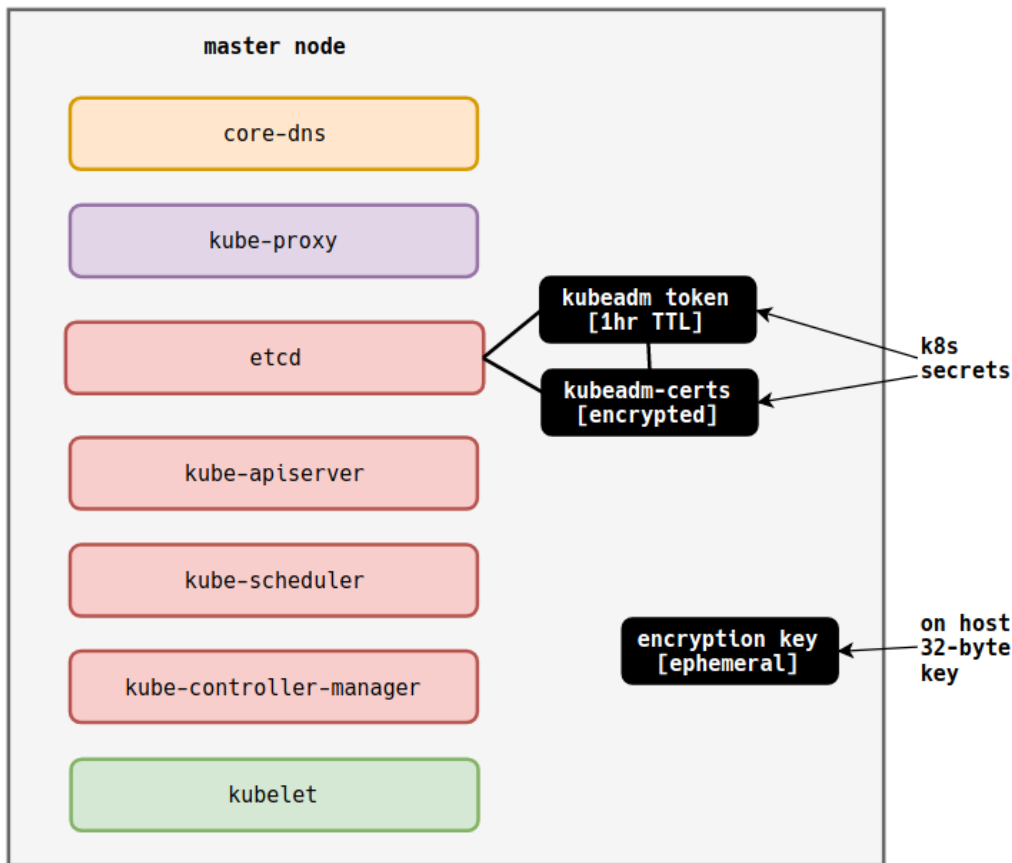
Conventionally, after installing a CNI plugin, users copy PKI information across 2 more master nodes and run a `kubeadm` command to add new control plane nodes. This results in a 3 node control plane.

1.14 introduced the `--upload-certs` flag to the `kubeadm init` command. As detailed in the [KEP](#), it has the following impact.

1. Creates an encryption key on the host.
2. Encrypts certificates and keys with the encryption key.
3. Adds encrypted data to the `kubeadm-certs` secret in the `kube-system` namespace.
4. Links the `kubeadm-certs` secret to a `kubeadm` token with a **1 hour TTL**.

Taking the example above, if we run the following command, it will result in the new diagram below.

```
kubeadm init --experimental-upload-certs
```



When the kubeadm token expires, so does the kubeadm-certs secret. Also, whenever the init phase upload-certs is run, a new encryption key is created. Ensuring that if kube-apiserver is compromised during the adding of a master node, the secret is encrypted and meaningless to the attacker. This example demonstrates running --experimental-upload-certs during cluster bootstrap. It is also possible to tap into phases, using kubeadm init phase upload-certs to achieve the above on an existing master. This will be detailed in the walkthrough below.

Lastly, it is important to note the token generated during upload-certs is only a proxy used to bind a TTL to the kubeadm-cert secret. You still need a conventional kubeadm token to join a control plane. This is the same token you would need to join a worker.

To add another control plane (master) node, a user can run the following command.

```
kubeadm join ${API_SERVER_PROXY_IP}:${API_SERVER_PROXY_PORT} \
  --experimental-control-plane \
  --certificate-key=${ENCRYPTION_KEY} \
  --token ${KUBEADM_TOKEN} \
  --discovery-token-ca-cert-hash ${APISERVER_CA_CERT_HASH}
```

## Walkthrough: Creating the HA Control Plane

This walkthrough will guide you to creating a new Kubernetes cluster with a 3 node control plane. It will demonstrate joining a control plane right after bootstrap and how to add another control plane node after the bootstrap tokens have expired.

1. Create 4 hosts (vms, baremetal, etc).

These hosts will be referred to as loadbalancer, master0, master1, and master2.

master hosts may need 2 vCPUs and 2GB of RAM available. You can get around these requirements during testing by ignoring pre-

flight checks. See the kubeadm documentation for more details.

2. Record the host IPs for later use.

## Setup the Load Balancer

In this section, we'll run a simple NGINX load balancer to provide a single endpoint for our control plane. This load balancer example is not meant for production scenarios.

1. SSH to the loadbalancer host.
2. Create the directory `/etc/nginx`.

```
mkdir /etc/nginx
```

3. Add and edit the file `/etc/nginx/nginx.conf`.

```
vim /etc/nginx/nginx.conf
```

4. Inside the file, add the following configuration.

```
events { }

stream {
    upstream stream_backend {
        least_conn;
        # REPLACE WITH master0 IP
        server 192.168.122.160:6443;
        # REPLACE WITH master1 IP
        server 192.168.122.161:6443;
        # REPLACE WITH master2 IP
        server 192.168.122.162:6443;
    }

    server {
        listen        6443;
        proxy_pass     stream_backend;
        proxy_timeout  3s;
        proxy_connect_timeout 1s;
    }
}
```

5. Alter each line above with a `REPLACE` comment above it.
6. Start NGINX.

```
docker run --name proxy \
-v /etc/nginx/nginx.conf:/etc/nginx/nginx.conf:ro \
-p 6443:6443 \
-d nginx
```

7. Verify you can reach NGINX at its address.

```
curl 192.168.122.170
```

output:

```
curl: (52) Empty reply from server
```

## Install Kubernetes Binaries on Master Hosts

1. Complete all pre-requisites and installation on master nodes.
  - a. See: <https://kubernetes.io/docs/setup/independent/install-kubeadm>.

## Initialize the Cluster and Examine Certificates

1. SSH to the master0 host.
2. Create the directory `/etc/kubernetes/kubeadm`

```
mkdir /etc/kubernetes/kubeadm
```

3. Create and edit the file `/etc/kubernetes/kubeadm/kubeadm-config.yaml`.

```
vim /etc/kubernetes/kubeadm/kubeadm-config.yaml
```

4. Add the following configuration.

```
apiVersion: kubeadm.k8s.io/v1beta1
kind: ClusterConfiguration
kubernetesVersion: stable
# REPLACE with `loadbalancer` IP
controlPlaneEndpoint: "192.168.122.170:6443"
networking:
  podSubnet: 192.168.0.0/18
```

5. Alter the line with a REPLACE comment above it.
6. Initialize the cluster with `upload-certs` and `config` specified.

```
kubeadm init \
  --config=/etc/kubernetes/kubeadm/kubeadm-config.yaml \
  --experimental-upload-certs
```

7. Record the output regarding joining control plane nodes for later use.

output:

You can now join any number of the control-plane node running the following command on each as root:

```
kubeadm join 192.168.122.170:6443 --token nmiqmn.yls76lcyxg2wt36c \
--discovery-token-ca-cert-hash
sha256:5efac16c86e5f2ed6b20c6dbcbf3a9daa5bf75aa604097dbf49fdc3d1fd5ff7d \
--experimental-control-plane --certificate-key
828fc83b950fca2c3bda129bcd0a4ffcd202cfb1a30b36abb901de1a3626a9df
```

Note the `certificate-key` that enables decrypting the kubeadm certs secret.

8. As your user, run the recommended kubeconfig commands for `kubectl` access.

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

9. Examine the `kubeadm-cert` secret in the `kube-system` namespace.

```
kubect1 get secrets -n kube-system kubeadm-certs -o yaml
```

You should see certs for etcd, kube-apiserver, and service accounts.

#### 10. Under ownerReferences, examine the name.

```
name: bootstrap-token-cwb9ra
```

This correlates to an ephemeral kubeadm token. When that token expires, so does this secret.

#### 11. List available tokens with kubeadm.

```
kubeadm token list
```

output:

TOKEN	TTL	EXPIRES	USAGES
DESCRIPTION			
cwb9ra.gegoj2eqddaf3yps	1h	2019-03-26T19:38:18Z	<none>
Proxy for managing TTL for the kubeadm-certs secret			
nmiqmn.yls76lcyxg2wt36c	23h	2019-03-27T17:38:18Z	
authentication, signing	<none>		

Note that `cwb9ra` is the owner reference in the above step. This is **not** a join token, instead a proxy that enables ttl on `kubeadm-certs`. We still need to use the `nmiqmn` token when joining.

#### 12. Install calico CNI-plugin with a pod CIDR matching the podSubnet configured above.

```
kubect1 apply -f
https://gist.github.com/joshrosso/ed1f5ea5a2f47d86f536e9eee3f1a2c2/raw/dfd95b9230fb3f75543706f3a95989964f36b154/calico-3.5.yaml
```

#### 13. Verify 1 node is Ready.

```
kubect1 get nodes
```

output:

NAME	STATUS	ROLES	AGE	VERSION
192-168-122-160	Ready	master	79m	v1.14.0

#### 14. Verify kube-system pods are Running.

```
kubect1 get pods -n kube-system
```

output:

NAME	READY	STATUS	RESTARTS	AGE
calico-node-mphpw	1/1	Running	0	58m
coredns-fb8b8dccf-c6s9q	1/1	Running	0	80m
coredns-fb8b8dccf-mxzrm	1/1	Running	0	80m
etcd-192-168-122-160	1/1	Running	0	79m
kube-apiserver-192-168-122-160	1/1	Running	0	79m

kube-controller-manager-192-168-122-160	1/1	Running	0	79m
kube-proxy-dpxhx	1/1	Running	0	80m
kube-scheduler-192-168-122-160	1/1	Running	0	79m

## Add the Second Master

1. SSH to the `master1` host.
2. Run the recorded join command from the previous section.

```
kubeadm join 192.168.122.170:6443 --token nmiqmn.yls76lcyxg2wt36c \
--discovery-token-ca-cert-hash
sha256:5efac16c86e5f2ed6b20c6dbcbf3a9daa5bf75aa604097dbf49fdc3d1fd5ff7d \
--experimental-control-plane \
--certificate-key
828fc83b950fca2c3bda129bcd0a4ffcd202cfb1a30b36abb901de1a3626a9df
```

3. After completion, verify there are now 2 nodes.

```
kubectl get nodes
```

output:

NAME	STATUS	ROLES	AGE	VERSION
192-168-122-160	Ready	master	22m	v1.14.0
192-168-122-161	Ready	master	34s	v1.14.0

4. Verify new pods have been created.

```
kubectl get pods -n kube-system
```

output:

NAME	READY	STATUS	RESTARTS	AGE
calico-node-cq5nt	1/1	Running	0	60s
calico-node-sp5w	1/1	Running	0	13m
coredns-fb8b8dccc-r9sc8	1/1	Running	0	22m
coredns-fb8b8dccc-wlcm4	1/1	Running	0	22m
etcd-192-168-122-160	1/1	Running	0	21m
etcd-192-168-122-161	1/1	Running	0	59s
kube-apiserver-192-168-122-160	1/1	Running	0	21m
kube-apiserver-192-168-122-161	1/1	Running	0	59s
kube-controller-manager-192-168-122-160	1/1	Running	0	21m
kube-controller-manager-192-168-122-161	1/1	Running	0	60s
kube-proxy-tflhf	1/1	Running	0	60s
kube-proxy-vthjr	1/1	Running	0	22m
kube-scheduler-192-168-122-160	1/1	Running	0	22m
kube-scheduler-192-168-122-161	1/1	Running	0	59s

## Add the Third Master with New Tokens

This section joins the third and final master. However, we will first delete all existing `kubeadm` tokens. This approach demonstrates how you could add masters when the Kubernetes cluster is already running.

1. On an existing master, list all tokens.

```
kubeadm token list
```

2. Delete all existing tokens.

```
kubeadm token delete cwb9ra.gegoj2eqddaf3yps
kubeadm token delete nmiqmn.yls76lcyxg2wt36c
```

Now the previously recorded join command will not work as the kubeadm-certs secret has expired and been deleted, the encryption key is no longer valid, and the join token will not work.

### 3. Create a new token with a 10 minute TTL.

```
kubeadm token create --ttl 10m --print-join-command
```

output:

```
kubeadm join 192.168.122.170:6443 \
  --token xaw58o.0fjg0xp0ohpucwhr \
  --discovery-token-ca-cert-hash
sha256:5efac16c86e5f2ed6b20c6dbcbf3a9daa5bf75aa604097dbf49fdc3d1fd5ff7d
```

Note the IP above must reflect your loadbalancer host.

### 4. Run the upload-certs phase of kubeadm init.

```
kubeadm init phase upload-certs --experimental-upload-certs
```

output:

```
[upload-certs] Storing the certificates in ConfigMap "kubeadm-certs" in
the "kube-system" Namespace
[upload-certs] Using certificate key:
9555b74008f24687eb964bd90a164ecb5760a89481d9c55a77c129b7db438168
```

### 5. SSH to the master2 host.

### 6. Use the outputs from the previous steps to run a control-plane join command.

```
kubeadm join 192.168.122.170:6443 \
  --experimental-control-plane \
  --certificate-key
9555b74008f24687eb964bd90a164ecb5760a89481d9c55a77c129b7db438168 \
  --token xaw58o.0fjg0xp0ohpucwhr \
  --discovery-token-ca-cert-hash
sha256:5efac16c86e5f2ed6b20c6dbcbf3a9daa5bf75aa604097dbf49fdc3d1fd5ff7d
```

### 7. After completion, verify there are now 3 nodes.

```
kubect1 get nodes
```

output:

NAME	STATUS	ROLES	AGE	VERSION
192-168-122-160	Ready	master	50m	v1.14.0
192-168-122-161	Ready	master	28m	v1.14.0
192-168-122-162	Ready	master	3m30s	v1.14.0

## Summary